Stock Market Project – Kafka and AWS

Create an EC2 instance in AWS:

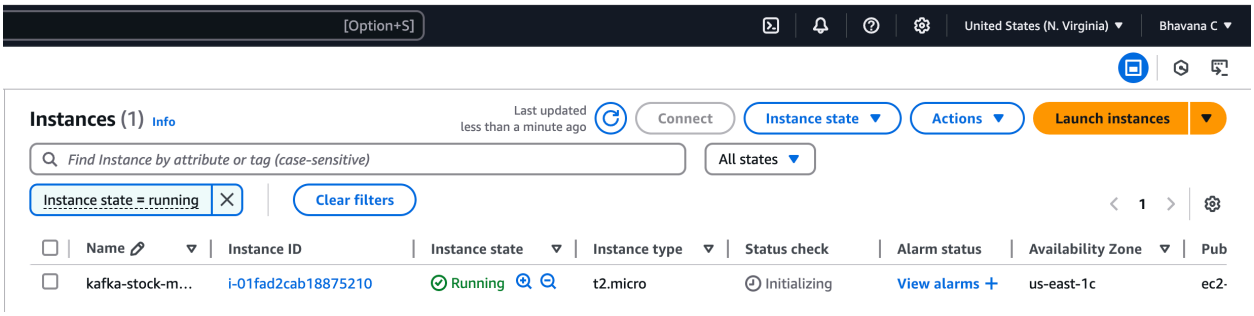

Start EC2:
cd "/Users/bhavanachinnamgari/Projects/Kafka_project_stocks"
ssh -i "kafka-stock-market-key.pem" ec2-user@ec2-75-101-243-117.compute-1.amazonaws.com

```
(base) bhavanachinnamgari@Bhavanas-Laptop Kafka_project_stocks % chmod 400 kafka-stock-market-key.pem
(base) bhavanachinnamgari@Bhavanas-Laptop Kafka_project_stocks % ssh -i "kafka-stock-market-key.pem" ec2-user@ec2-75-101-243-117.compute-1.amazonaws.com
       __|  __|_  )
       _|  (     /   Amazon Linux 2
      ___|\___|___|
https://aws.amazon.com/linux/amazon-linux-2
[ec2-user@ip-172-31-11-234 ~]$
```

Download Kafka and unzip the folder:
wget https://dlcdn.apache.org/kafka/3.9.0/kafka-3.9.0-src.tgz
tar -xvf kafka_2.12-3.9.0.tgz

```
[ec2-user@ip-172-31-11-234 ~]$ wget https://dlcdn.apache.org/kafka/3.9.0/kafka_2.12-3.9.0.tgz
--2025-03-09 02:18:45--  https://dlcdn.apache.org/kafka/3.9.0/kafka_2.12-3.9.0.tgz
Resolving dlcdn.apache.org (dlcdn.apache.org)... 151.101.2.132, 2a04:4e42::644
Connecting to dlcdn.apache.org (dlcdn.apache.org)|151.101.2.132|:443... connected.
HTTP request sent, awaiting response... 200 OK
Length: 122204110 (117M) [application/x-gzip]
Saving to: 'kafka_2.12-3.9.0.tgz'

100%[=========================================================================================================

2025-03-09 02:18:50 (108 MB/s) - 'kafka_2.12-3.9.0.tgz' saved [122204110/122204110]

[ec2-user@ip-172-31-11-234 ~]$ ls
kafka_2.12-3.9.0.tgz
[ec2-user@ip-172-31-11-234 ~]$ tar -xvf kafka_2.12-3.9.0.tgz
kafka_2.12-3.9.0/
kafka_2.12-3.9.0/LICENSE
kafka_2.12-3.9.0/NOTICE
kafka_2.12-3.9.0/bin/
kafka_2.12-3.9.0/bin/kafka-verifiable-producer.sh
```

Check for Java JDK, install if not present: sudo yum install java-1.8.0-openjdk
Change path into folder: cd kafka_2.12-3.9.0
Run zookeeper in Terminal-1: bin/zookeeper-server-start.sh config/zookeeper.properties

---

Start EC2 in Terminal-2 and run Kafka server:
export KAFKA_HEAP_OPTS="-Xmx256M -Xms128M"
cd kafka_2.12-3.9.0
bin/kafka-server-start.sh config/server.properties

```
[ec2-user@ip-172-31-11-234 ~]$ export KAFKA_HEAP_OPTS="-Xmx256M -Xms128M"
[ec2-user@ip-172-31-11-234 ~]$ cd kafka_2.12-3.9.0
[ec2-user@ip-172-31-11-234 kafka_2.12-3.9.0]$ bin/kafka-server-start.sh config/server.properties
OpenJDK 64-Bit Server VM warning: If the number of processors is expected to increase from one,
[2025-03-09 02:30:13,414] INFO Registered kafka:type=kafka.Log4jController MBean (kafka.utils.Log
[2025-03-09 02:30:13,976] INFO Setting -D jdk.tls.rejectClientInitiatedRenegotiation=true to disa
[2025-03-09 02:30:14,162] INFO Registered signal handlers for TERM, INT, HUP (org.apache.kafka.co
[2025-03-09 02:30:14,165] INFO starting (kafka.server.KafkaServer)
[2025-03-09 02:30:14,167] INFO Connecting to zookeeper on localhost:2181 (kafka.server.KafkaServe
[2025-03-09 02:30:14,196] INFO [ZooKeeperClient Kafka server] Initializing a new session to local
[2025-03-09 02:30:14,203] INFO Client environment:zookeeper.version=3.8.4-9316c2a7a97e1666d8f459
```

Change properties file with EC2 port number: sudo nano config/server.properties

```
# Listener name, hostname and port the broker will advertise to clients.
# If not set, it uses the value for "listeners".
advertised.listeners=PLAINTEXT://75.101.243.117:9092
```

CTRL + X, Y, RETURN

Run both Zookeeper and Kafka again and check below in Kafka server:

```
[2025-03-09 02:36:25,893] INFO KafkaConfig values:
        advertised.listeners = PLAINTEXT://75.101.243.117:9092
```

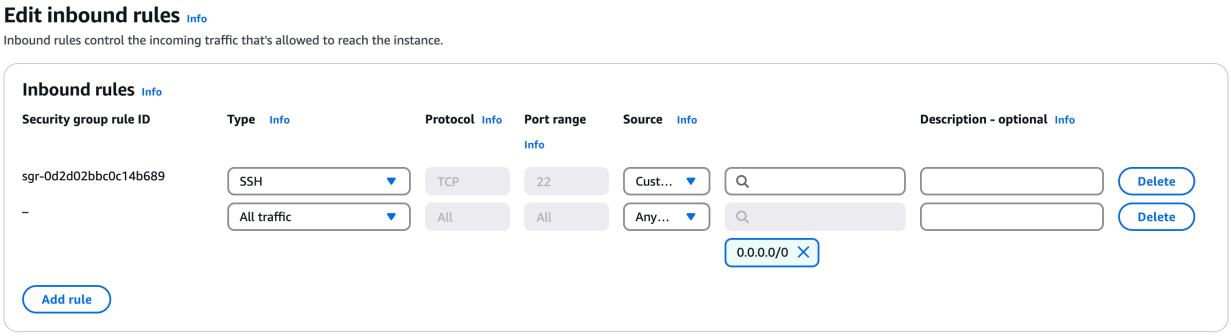

Now, we have Zookeeper and Kafka setup done!

Let’s create Topic, Producer and Consumer. Start EC2 instances in new terminals and Goto Kafka folder

Create the topic with replication factor and partitions: Terminal-3

bin/kafka-topics.sh --create --topic stocks-test6 --bootstrap-server 75.101.243.117:9092 --partitions 1 --replication-factor 1

General command: bin/kafka-topics.sh --create --topic {topic name} --bootstrap-server {EC2 instance/localhost}:9092 --partitions 1 --replication-factor 1

“stocks-test6” Topic gets created.

Start Producer: Terminal-3

bin/kafka-console-producer.sh --topic stocks-test6 --bootstrap-server 75.101.243.117:9092

Start Consumer in a new Terminal-4

cd kafka_2.12-3.9.0

bin/kafka-console-consumer.sh --topic stocks-test6 --bootstrap-server 75.101.243.117:9092 --from-beginning

---

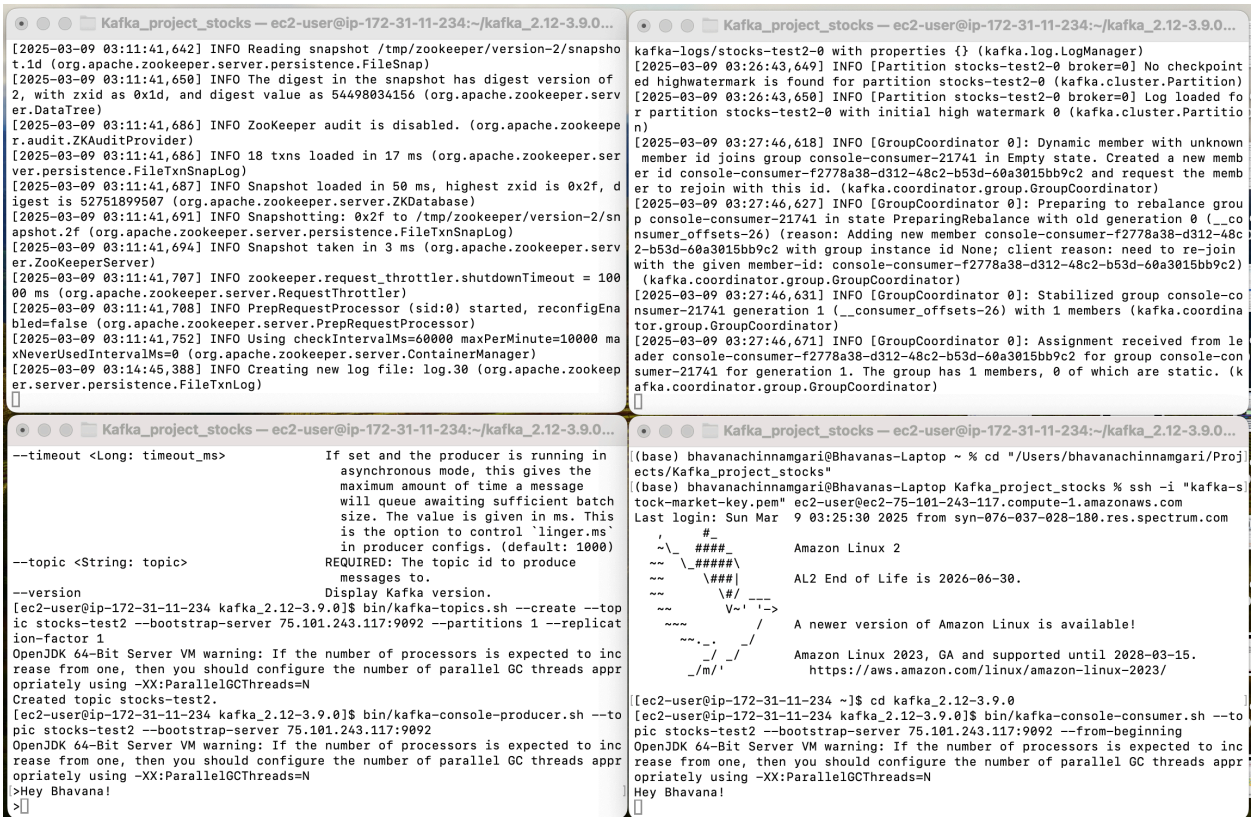

Producer and Consumer are working well!

Login to AWS using CLI with Access ID and Secret ID. Write the python code for the kafka project to display output in Jupyter/ VS code for testing. Then use stocks_data.csv to populate them in S3 bucket.

Create S3 bucket to save the output:

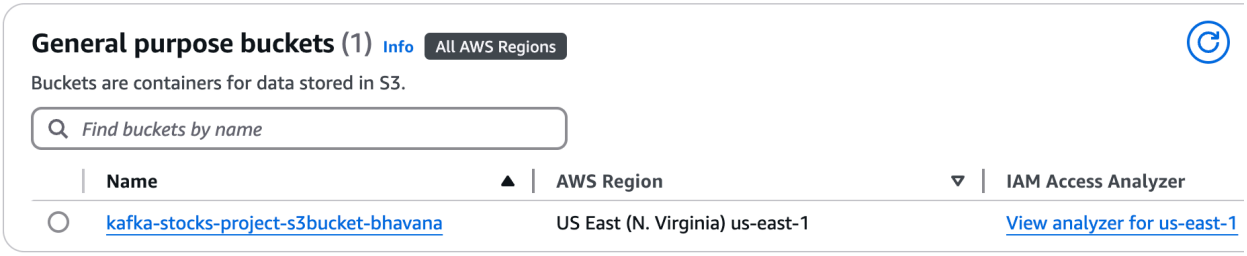

Run for 5 sec and check the console if the json data is being populated.

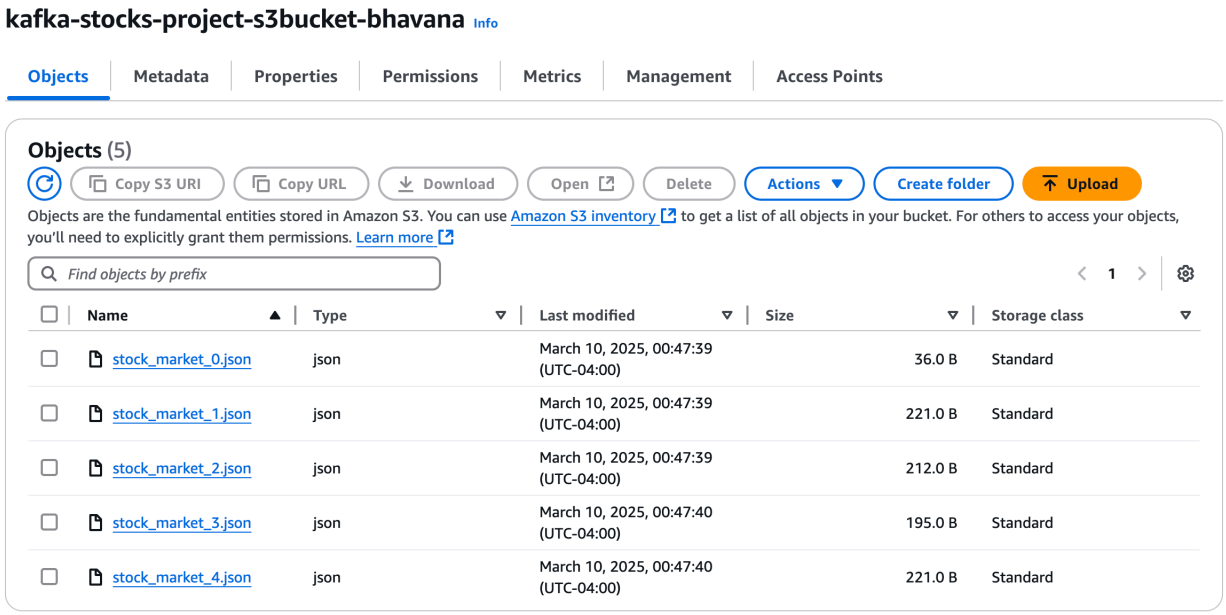

Create a AWS Crawler using our S3 bucket

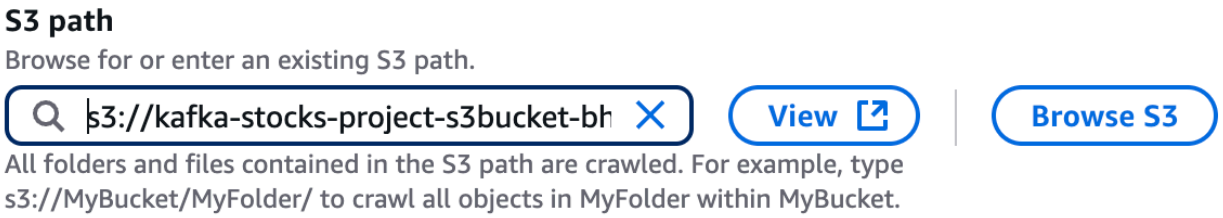

Create a Role to give access to AWS Glue:

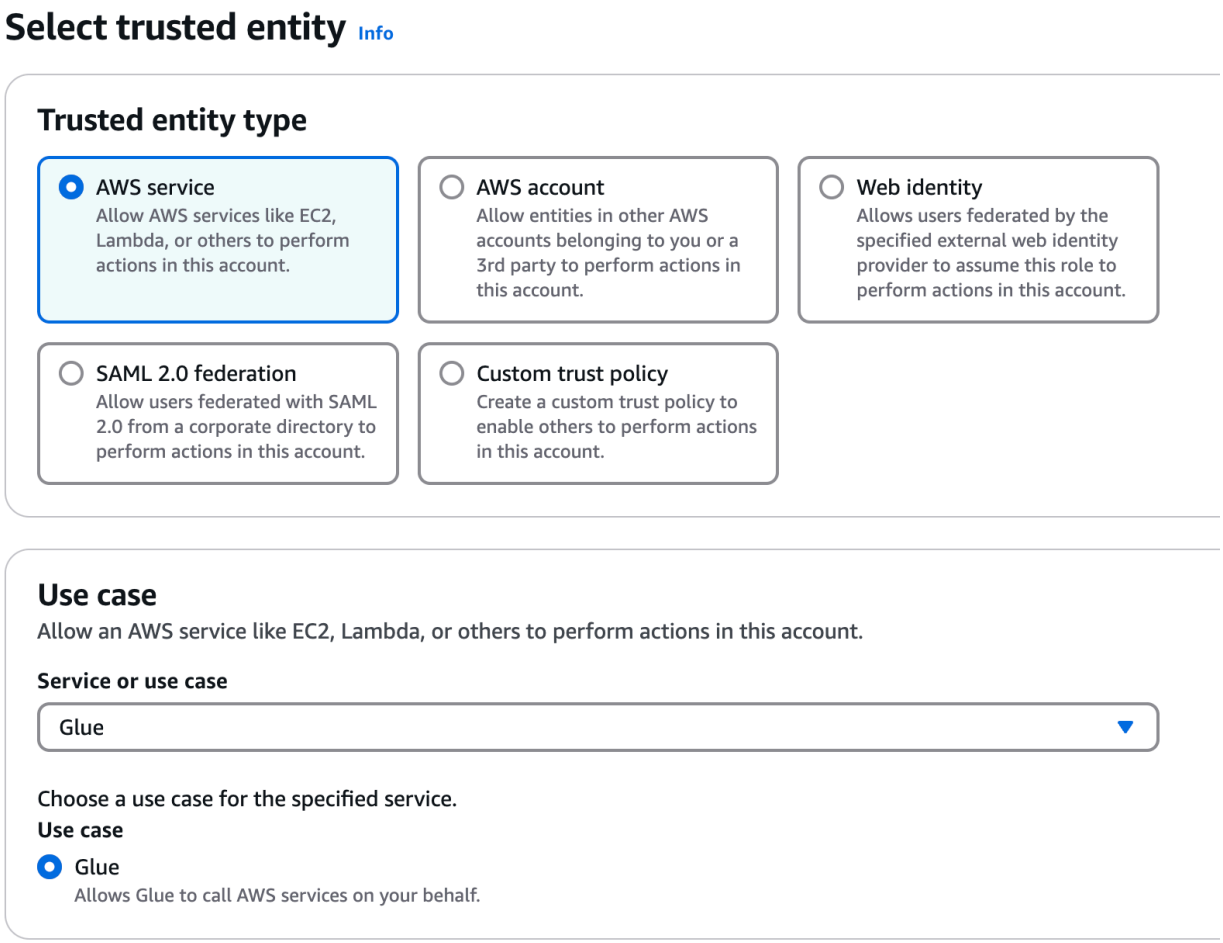

Create a Database “stock-market-db” to store the table created by Crawler:

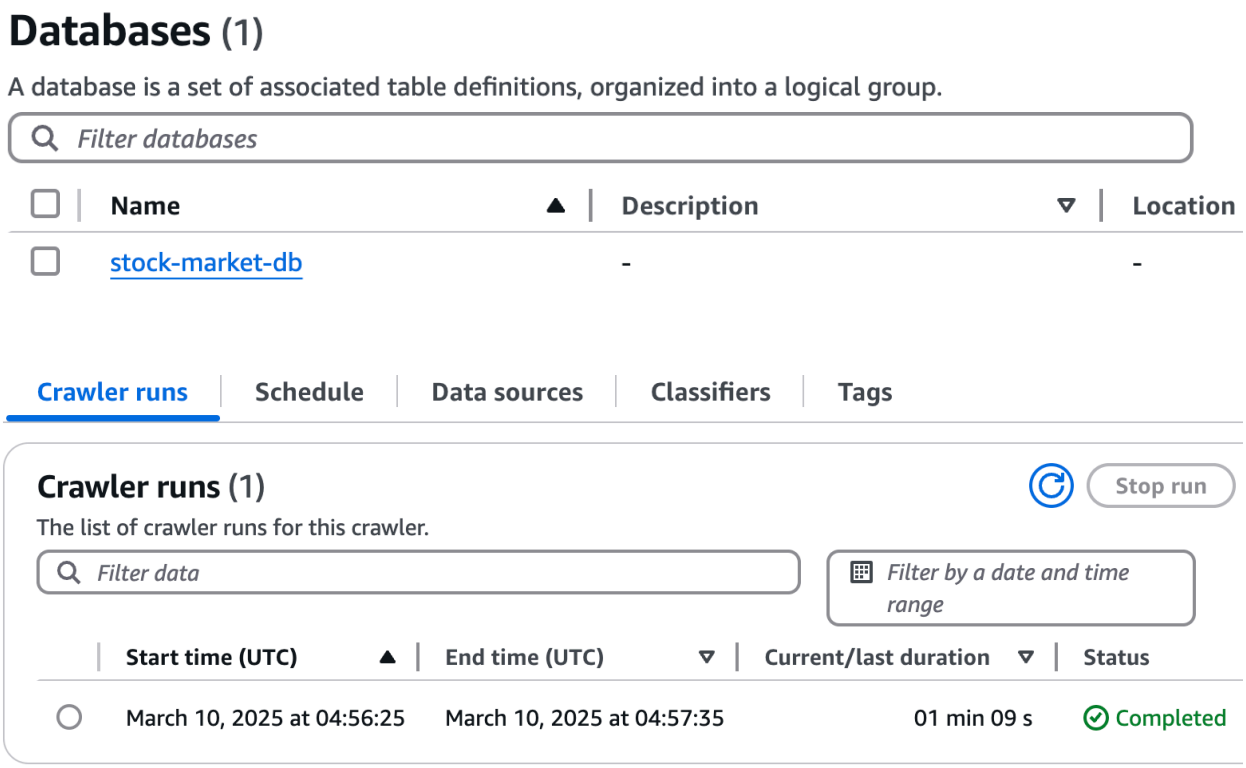

A Table is created under the database “stock-market-db”. Now we can execute queries using AWS Athena. Create another S3 bucket to save the query information.